

BDA LAB REPORT

Name – Shivam Raj

USN – 1BM17CS095

BATCH – B3

Q1.)

Perform the following DB operations using MongoDB.

- A. Create a database “Student” with the following attributes Rollno, Age, ContactNo, Email-Id.
- B. Insert appropriate values.
- C. Write query to update Email-Id of a student with rollno 10.
- D. Replace the student name from “ABC” to “FEM” of rollno 11.
- E. Export the created table into local file system.
- F. Drop the table.
- G. Import a given csv dataset from local file system into mongodb collection.

CODE – 1

```
> use student;
```

```
> db;
```

```
> show dbs;
```

```
> db.createCollection("Stud");
```

```
> db.Stud.insert({_id:1, name : "Shivam", rollno:1, age:21, ContactNo:999, EmailId: "shiv@r.com"});
```

```
> db.Stud.insert({_id:2, name : "Shiv", rollno:10, age:21, ContactNo:9999, EmailId:
"aBCD@b.com"});

> db.Stud.insert({_id:3, name : "ABC", rollno:11, age:21, ContactNo:999999, EmailId:
"aCD@b.com"});

> db.Stud.insert({_id:4, name : "ADAABC", rollno:12, age:22, ContactNo:99999988, EmailId:
"DNEW@b.com"});
```

```
> db.Stud.find();

> db.Stud.update({rollno:10},{ $set:{EmailId:"hello@b.com"}});

> db.Stud.find({rollno:10});

> db.Stud.replaceOne({"rollno":11},{ "name": "FEM", "rollno":11, "age":21,
"ContactNo":9999, "EmailId": "aBCD@b.com"});

> exit
```

```
mongoexport -d student -c Stud -f name,rollno,age,ContactNo,EmailId --type=csv -o
F:\Mongo\Student.csv
```

```
> use student;

> db;

> db.Stud.drop();

> db.Stud.find();

> exit
```

```
mongoimport -d student -c Stud --type csv -f name,rollno,age,ContactNo,EmailId --file
F:\Mongo\Student.csv
```

```
> use student;

> db.Stud.find();
```

Write-up-1

Shivom Raj
LBMI7C S095

BDA Lab
B3

Lab-3 (I)

classmate

Date

Page

```
use student;
db;
db.createCollection("stud");
db.stud.insert({_id:1, name:"Shivom",
rollno:1, age:21, contactno:9999,
EmailId:"Shivom@x.com"});

db.stud.update({rollno:10}, {$set:{EmailId:
"hello@s.com"}});

db.stud.replaceOne({rollno:11}, {"name":"FERM",
"rollno":11, "age":21, "ContactNo":999,
"EmailId":"abc@s.com"});

mongoexport -d student -c stud -f name, rollno, age,
ContactNo, EmailId --type=csv -o D:\mongo\student.csv

db.stud.drop();
db.stud.find();

mongoimport -d student -c stud -f type csv -f
name, rollno, age, ContactNo, EmailId --file D:\mongo\
student.csv

db.stud.find();
```

Screenshot – 1

```
Command Prompt - mongo
> use student;
switched to db student
> db;
student
> show dbs;
admin    0.000GB
company  0.000GB
config   0.000GB
local    0.000GB
> db.createCollection("stud");
{ "ok" : 1 }
> db.stud.insert({_id:1,name:"Shivam",rollno:1,age:21,ContactNo:999, EmailId: "shiv@r.com"});
writeResult({ "nInserted" : 1 })
> db.stud.insert({_id:2, name: "Shiv", rollno:10, age:21, ContactNo:9999, EmailId: "aBCD@b.com"});
writeResult({ "nInserted" : 1 })
> db.stud.insert({_id:3, name: "ABC", rollno:11, age:21, ContactNo:999999, EmailId: "aCD@b.com"});
writeResult({ "nInserted" : 1 })
> db.stud.insert({_id:4, name: "ADAAABC", rollno:12, age:22, ContactNo:99999988, EmailId: "DNE@b.com"});
writeResult({ "nInserted" : 1 })
>
> db.stud.find();
{ "_id" : 1, "name" : "Shivam", "rollno" : 1, "age" : 21, "ContactNo" : 999, "EmailId" : "shiv@r.com" }
> db.stud.find();
{ "_id" : 2, "name" : "Shiv", "rollno" : 10, "age" : 21, "ContactNo" : 9999, "EmailId" : "aBCD@b.com" }
{ "_id" : 3, "name" : "ABC", "rollno" : 11, "age" : 21, "ContactNo" : 999999, "EmailId" : "aCD@b.com" }
{ "_id" : 4, "name" : "ADAAABC", "rollno" : 12, "age" : 22, "ContactNo" : 99999988, "EmailId" : "DNE@b.com" }
> db.stud.find();
{ "_id" : 2, "name" : "Shiv", "rollno" : 10, "age" : 21, "ContactNo" : 9999, "EmailId" : "aBCD@b.com" }
{ "_id" : 3, "name" : "ABC", "rollno" : 11, "age" : 21, "ContactNo" : 999999, "EmailId" : "aCD@b.com" }
{ "_id" : 4, "name" : "ADAAABC", "rollno" : 12, "age" : 22, "ContactNo" : 99999988, "EmailId" : "DNE@b.com" }
>
> db.stud.update(rollno:10,{set:{EmailId:"hello@b.com"}});
writeResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
> db.stud.find(rollno:10);
{ "_id" : 2, "name" : "Shiv", "rollno" : 10, "age" : 21, "ContactNo" : 9999, "EmailId" : "hello@b.com" }
> db.stud.replaceOne(rollno:11,{"name": "FEH", "rollno":11, "age":21, "ContactNo":9999, "EmailId": "aBCD@b.com"});
2020-10-08T15:58:11.369+0530 E QUERY [js] uncaught exception: SyntaxError: missing ) after argument list :
@shell:1:32
> db.stud.replaceOne(rollno:11,{"name": "FEH", "rollno":11, "age":21, "ContactNo":9999, "EmailId": "aBCD@b.com"});
{ "acknowledged" : true, "matchedCount" : 1, "modifiedCount" : 1 }
> exit
bye

C:\Program Files\MongoDB\Server\4.2\bin>mongoexport -d student -c Stud -f name,rollno,age,ContactNo,EmailId --type=csv -o D:\Wongo\Student.csv
2020-10-08T18:16:15.966+0530 connected to: mongod://localhost/
2020-10-08T18:16:16.318+0530 exported 3 records

C:\Program Files\MongoDB\Server\4.2\bin>mongo
MongoDB shell version v4.2.2

C:\Program Files\MongoDB\Server\4.2\bin>mongoexport -d student -c Stud -f name,rollno,age,ContactNo,EmailId --type=csv -o D:\Wongo\Student.csv
2020-10-08T18:16:15.966+0530 connected to: mongod://localhost/
2020-10-08T18:16:16.318+0530 exported 3 records

C:\Program Files\MongoDB\Server\4.2\bin>mongo
MongoDB shell version v4.2.2
connecting to: mongod://127.0.0.1:27017/?compressors=disabled&gssapiServiceName=mongod
Implicit session: session { "id" : UUID("14d8eb81-7163-431e-9db0-3593d68d3e4e") }
MongoDB server version: 4.2.2
Server has startup warnings:
2020-10-01T16:00:08.536+0530 I CONTROL [initandlisten]
2020-10-01T16:00:08.536+0530 I CONTROL [initandlisten] ** WARNING: Access control is not enabled for the database.
2020-10-01T16:00:08.537+0530 I CONTROL [initandlisten] ** Read and write access to data and configuration is unrestricted.
2020-10-01T16:00:08.586+0530 I CONTROL [initandlisten]

Enable MongoDB's free cloud-based monitoring service, which will then receive and display
metrics about your deployment (disk utilization, CPU, operation statistics, etc).

The monitoring data will be available on a MongoDB website with a unique URL accessible to you
and anyone you share the URL with. MongoDB may use this information to make product
improvements and to suggest MongoDB products and deployment options to you.

To enable free monitoring, run the following command: db.enableFreeMonitoring()
To permanently disable this reminder, run the following command: db.disableFreeMonitoring()
...

> use student;
switched to db student
> db;
student
> db.stud.find();
{ "_id" : 1, "name" : "Shivam", "rollno" : 1, "age" : 21, "ContactNo" : 999, "EmailId" : "shiv@r.com" }
> db.stud.find();
{ "_id" : 2, "name" : "Shiv", "rollno" : 10, "age" : 21, "ContactNo" : 9999, "EmailId" : "hello@b.com" }
{ "_id" : 3, "name" : "FEH", "rollno" : 11, "age" : 21, "ContactNo" : 9999, "EmailId" : "aBCD@b.com" }
{ "_id" : 4, "name" : "ADAAABC", "rollno" : 12, "age" : 22, "ContactNo" : 99999988, "EmailId" : "DNE@b.com" }
> db.stud.drop();
true
> db.stud.find();
>
> db.stud.find();
{ "_id" : 1, "name" : "Shivam", "rollno" : 1, "age" : 21, "ContactNo" : 999, "EmailId" : "shiv@r.com" }
>
> db.stud.find();
>
> exit
bye

C:\Program Files\MongoDB\Server\4.2\bin>mongoimport -d student -c Stud --type csv -f name,rollno,age,ContactNo,EmailId --file D:\Wongo\Student.csv
2020-10-08T18:31:58.871+0530 connected to: mongod://localhost/
```

```

Command Prompt - mongo
>
> db.Stud.find();
> exit
bye

C:\Program Files\MongoDB\Server\4.2\bin>mongoimport -d student -c Stud --type csv -f name,rollno,age,ContactNo,EmailId --file D:\Mongo\Student.csv
2020-10-08T18:31:58.871+0530   connected to: mongodb://localhost/
2020-10-08T18:31:59.251+0530   4 document(s) imported successfully. 0 document(s) failed to import.

C:\Program Files\MongoDB\Server\4.2\bin>mongo
MongoDB shell version v4.2.2
connecting to: mongodb://127.0.0.1:27017/?compressors=disabled&gssapiServiceName=mongodb
Implicit session: session {"id": "UUID("669d582b-cc6f-4daa-a1bf-ebbdfc1a0e48")"}
MongoDB server version: 4.2.2
Server has startup warnings:
2020-10-01T16:00:08.536+0530 I CONTROL [initandlisten]
2020-10-01T16:00:08.536+0530 I CONTROL [initandlisten] ** WARNING: Access control is not enabled for the database.
2020-10-01T16:00:08.537+0530 I CONTROL [initandlisten] **           Read and write access to data and configuration is unrestricted.
2020-10-01T16:00:08.586+0530 I CONTROL [initandlisten]
***
Enable MongoDB's free cloud-based monitoring service, which will then receive and display
metrics about your deployment (disk utilization, CPU, operation statistics, etc).

The monitoring data will be available on a MongoDB website with a unique URL accessible to you
and anyone you share the URL with. MongoDB may use this information to make product
improvements and to suggest MongoDB products and deployment options to you.

To enable free monitoring, run the following command: db.enableFreeMonitoring()
To permanently disable this reminder, run the following command: db.disableFreeMonitoring()
***

> use student;
switched to db student
> db.Stud.find();
{ "_id" : ObjectId("5f7fd8c65640bba1b0354304"), "name" : "name", "rollno" : "rollno", "age" : "age", "ContactNo" : "ContactNo", "EmailId" : "EmailId" }
{ "_id" : ObjectId("5f7fd8c65640bba1b0354305"), "name" : "shiv", "rollno" : 10, "age" : 21, "ContactNo" : 9999, "EmailId" : "hello@b.com" }
{ "_id" : ObjectId("5f7fd8c65640bba1b035430F"), "name" : "FEH", "rollno" : 11, "age" : 21, "ContactNo" : 9999, "EmailId" : "sBCD@b.com" }
{ "_id" : ObjectId("5f7fd8c65640bba1b035431B"), "name" : "ADABCB", "rollno" : 12, "age" : 22, "ContactNo" : 999999988, "EmailId" : "DNEA@b.com" }
> db.stud.find();
{ "_id" : 1, "name" : "Shivam", "rollno" : 1, "age" : 21, "ContactNo" : 999, "EmailId" : "shiv@r.com" }
>
>

```

Q2.)

Perform the following DB operations using MongoDB.

- A. Create a collection by name Customers with the following attributes.
Cust_id, Acc_Bal, Acc_Type.
- B. Insert at least 5 values into the table.
- C. Write a query to display those records whose total account balance is greater than 1200 of account type 'Z' for each customer_id.
- D. Determine Minimum and Maximum account balance for each customer_id.
- E. Export the created collection into local file system
- F. Drop the table.
- G. Import a given csv dataset from local file system into mongodb collection.

CODE – 2

```
> use mynew;
> db;
> show dbs;
> db.createCollection("Customer");
> db.Customer.insert({cust_id:1,Acc_bal:1500,Acc_type:"Z"});
> db.Customer.insert({cust_id:2,Acc_bal:3000,Acc_type:"A"});
> db.Customer.insert({cust_id:3,Acc_bal:1200,Acc_type:"A"});
> db.Customer.insert({cust_id:2,Acc_bal:1200,Acc_type:"Z"});
> db.Customer.insert({cust_id:3,Acc_bal:1280,Acc_type:"Z"});

> db.Customer.find();
> db.Customer.find({Acc_bal:{$gt:1200}, Acc_type:"Z"});
> db.Customer.aggregate([
... ..  {
... ..  $group: {
```

```
... ..    _id: "$cust_id",
... ..    min_bal: {$min: "$Acc_bal"},
... ..    max_bal: {$max: "$Acc_bal"}
... ..    }
... ..    }
... ..    ]);
```

```
>
```

```
> db.Customer.find();
```

```
> exit
```

```
mongoexport -d mynew -c Customer -f cust_id,Acc_bal,Acc_type --type=csv -o
F:\Mongo\Customer.csv
```

```
mongo
```

```
> use mynew;
```

```
> db.Customer.drop();
```

```
> db.Customer.find();
```

```
> exit
```

```
mongoimport -d mynew -c Customer --type csv -f cust_id,Acc_bal,Acc_type --file
F:\Mongo\Customer.csv
```

```
mongo
```

```
> use mynew;
```

```
> db.Customer.find();
```


Write-up-2

Shivam Raj
IBM17C8095

BDA Lab
B3

Lab-3 (II)

classmate

Date

Page

use my new;

db;

db.createCollection("Customer");

db.Customer.insert({cust-id:1, Acc-bal:1500,
Acc-type:"Z"});

db.Customer.find();

db.Customer.find({Acc-bal:{>1200},
Acc-type:"Z"});

db.Customer.aggregate([

{

_id:"\$cust-id",

min-bal:{>min:"\$Acc-bal"},

max-bal:{>max:"\$Acc-bal"}
}]

}

}

];

mongoexport -d mynew -c Customer -f cust-id,
Acc-bal, Acc-type --type csv -o D:\mongo\customer.csv

db.customer.drop()

db.customer.find()

mongoimport -d mynew -c Customer --type csv -f cust-id,
Acc-bal, Acc-type --file D:\mongo\customer.csv

db.customer.find();

Screenshot – 2

```
Microsoft Windows [Version 10.0.18363.1882]
(c) 2019 Microsoft Corporation. All rights reserved.

C:\Users\SHIVAM>cd C:\Program Files\MongoDB\Server\4.2\bin

C:\Program Files\MongoDB\Server\4.2\bin>mongo
MongoDB shell version v4.2.2
connecting to: mongodb://127.0.0.1:27017/?compressors=disabled&gssapiServiceName=mongodb
Implicit session: session { "id" : UUID("9fcd5c5-5643-4afe-9b45-261b64e9b9fe") }
MongoDB server version: 4.2.2
Server has startup warnings:
2020-10-01T16:00:08.536+0530 I CONTROL [initandlisten]
2020-10-01T16:00:08.536+0530 I CONTROL [initandlisten] ** WARNING: Access control is not enabled for the database.
2020-10-01T16:00:08.537+0530 I CONTROL [initandlisten] **           Read and write access to data and configuration is unrestricted.
2020-10-01T16:00:08.586+0530 I CONTROL [initandlisten]
---
Enable MongoDB's free cloud-based monitoring service, which will then receive and display
metrics about your deployment (disk utilization, CPU, operation statistics, etc).

The monitoring data will be available on a MongoDB website with a unique URL accessible to you
and anyone you share the URL with. MongoDB may use this information to make product
improvements and to suggest MongoDB products and deployment options to you.

To enable free monitoring, run the following command: db.enableFreeMonitoring()
To permanently disable this reminder, run the following command: db.disableFreeMonitoring()
---

> use mynew;
switched to db mynew
> db;
mynew
> show dbs;
admin      0.000GB
company    0.000GB
config     0.000GB
local      0.000GB
student    0.000GB
> db.createCollection("Customer");
{ "ok" : 1 }
> db.Customer.insert({cust_id:1,Acc_bal:2500,Acc_type:"Z"});
WriteResult({ "ninserted" : 1 })
> db.Customer.insert({cust_id:2,Acc_bal:3000,Acc_type:"A"});
WriteResult({ "ninserted" : 1 })
> db.Customer.insert({cust_id:3,Acc_bal:1200,Acc_type:"A"});
WriteResult({ "ninserted" : 1 })
> db.Customer.insert({cust_id:2,Acc_bal:2400,Acc_type:"A"});
WriteResult({ "ninserted" : 1 })
> db.Customer.insert({cust_id:4,Acc_bal:1200,Acc_type:"Z"});
WriteResult({ "ninserted" : 1 })
> db.Customer.find();
{ "_id" : ObjectId("5f7f231c0a608fe9f8550dc5"), "cust_id" : 1, "Acc_bal" : 2500, "Acc_type" : "Z" }
{ "_id" : ObjectId("5f7f2350a0a08fe9f8550dc6"), "cust_id" : 2, "Acc_bal" : 3000, "Acc_type" : "A" }
{ "_id" : ObjectId("5f7f236c0a08fe9f8550dc7"), "cust_id" : 3, "Acc_bal" : 1200, "Acc_type" : "A" }
{ "_id" : ObjectId("5f7f24660a08fe9f8550dc8"), "cust_id" : 2, "Acc_bal" : 2400, "Acc_type" : "A" }
{ "_id" : ObjectId("5f7f250f0a08fe9f8550dc9"), "cust_id" : 4, "Acc_bal" : 1200, "Acc_type" : "Z" }
>
> db.Customer.find({Acc_bal:{$gt:1200}, Acc_type:"Z"});
{ "_id" : ObjectId("5f7f231c0a608fe9f8550dc5"), "cust_id" : 1, "Acc_bal" : 2500, "Acc_type" : "Z" }
>
> db.Customer.aggregate([
... {
...   $group:
...   {
...     _id: "$cust_id",
...     min_bal: {$min: "$Acc_bal"},
...     max_bal: {$max: "$Acc_bal"}
...   }
... });
{ "_id" : 4, "min_bal" : 1200, "max_bal" : 1200 }
{ "_id" : 3, "min_bal" : 1200, "max_bal" : 1200 }
{ "_id" : 1, "min_bal" : 2500, "max_bal" : 2500 }
{ "_id" : 2, "min_bal" : 2400, "max_bal" : 3000 }
> db.Customer.find();
{ "_id" : ObjectId("5f7f231c0a608fe9f8550dc5"), "cust_id" : 1, "Acc_bal" : 2500, "Acc_type" : "Z" }
{ "_id" : ObjectId("5f7f2350a0a08fe9f8550dc6"), "cust_id" : 2, "Acc_bal" : 3000, "Acc_type" : "A" }
{ "_id" : ObjectId("5f7f236c0a08fe9f8550dc7"), "cust_id" : 3, "Acc_bal" : 1200, "Acc_type" : "A" }
{ "_id" : ObjectId("5f7f24660a08fe9f8550dc8"), "cust_id" : 2, "Acc_bal" : 2400, "Acc_type" : "A" }
{ "_id" : ObjectId("5f7f250f0a08fe9f8550dc9"), "cust_id" : 4, "Acc_bal" : 1200, "Acc_type" : "Z" }
> exit
bye

C:\Program Files\MongoDB\Server\4.2\bin>mongoexport -d mynew -c Customer -f cust_id,Acc_bal,Acc_type --type=csv -o D:\Wongo\Customer.csv
2020-10-08T20:23:42.276+0530   connected to: mongodb://localhost/
2020-10-08T20:23:42.321+0530   exported 5 records

C:\Program Files\MongoDB\Server\4.2\bin>mongo
MongoDB shell version v4.2.2
connecting to: mongodb://127.0.0.1:27017/?compressors=disabled&gssapiServiceName=mongodb
Implicit session: session { "id" : UUID("435b1079-f100-4993-a7d0-580cd492b0f7") }
MongoDB server version: 4.2.2
Server has startup warnings:
2020-10-01T16:00:08.536+0530 I CONTROL [initandlisten]
2020-10-01T16:00:08.536+0530 I CONTROL [initandlisten]
2020-10-01T16:00:08.537+0530 I CONTROL [initandlisten] ** WARNING: Access control is not enabled for the database.
2020-10-01T16:00:08.586+0530 I CONTROL [initandlisten] **           Read and write access to data and configuration is unrestricted.
2020-10-01T16:00:08.586+0530 I CONTROL [initandlisten]
---
```

```

C:\Program Files\MongoDB\Server\4.2\bin>mongoimport -d mynew -c Customer --type csv -f cust_id,Acc_bal,Acc_type --file D:\Mongo\Customer.csv
2020-10-08T20:26:31.320+0530   connected to: mongod://127.0.0.1:27017/?compressors=disabled&gssapiServiceName=mongod
2020-10-08T20:26:31.516+0530   6 document(s) imported successfully. 0 document(s) failed to import.

C:\Program Files\MongoDB\Server\4.2\bin>mongo
MongoDB shell version v4.2.2
connecting to: mongod://127.0.0.1:27017/?compressors=disabled&gssapiServiceName=mongod
Implicit session: session { "id": "UUID('87c4cb33-338c-4949-aace-b5012e1977c6') " }
MongoDB server version: 4.2.2
Server has startup warnings:
2020-10-01T16:00:08.536+0530 I CONTROL [initandlisten]
2020-10-01T16:00:08.536+0530 I CONTROL [initandlisten] ** WARNING: Access control is not enabled for the database.
2020-10-01T16:00:08.537+0530 I CONTROL [initandlisten] **           Read and write access to data and configuration is unrestricted.
2020-10-01T16:00:08.586+0530 I CONTROL [initandlisten]

Enable MongoDB's free cloud-based monitoring service, which will then receive and display
metrics about your deployment (disk utilization, CPU, operation statistics, etc).

The monitoring data will be available on a MongoDB website with a unique URL accessible to you
and anyone you share the URL with. MongoDB may use this information to make product
improvements and to suggest MongoDB products and deployment options to you.

To enable free monitoring, run the following command: db.enableFreeMonitoring()
To permanently disable this reminder, run the following command: db.disableFreeMonitoring()
---
> use mynew;
switched to db mynew
> db.Customer.drop();
true
> db.Customer.find();
[]
> exit
bye

C:\Program Files\MongoDB\Server\4.2\bin>mongoimport -d mynew -c Customer --type csv -f cust_id,Acc_bal,Acc_type --file D:\Mongo\Customer.csv
2020-10-08T20:26:31.320+0530   connected to: mongod://127.0.0.1:27017/?compressors=disabled&gssapiServiceName=mongod
2020-10-08T20:26:31.516+0530   6 document(s) imported successfully. 0 document(s) failed to import.

C:\Program Files\MongoDB\Server\4.2\bin>mongo
MongoDB shell version v4.2.2
connecting to: mongod://127.0.0.1:27017/?compressors=disabled&gssapiServiceName=mongod
Implicit session: session { "id": "UUID('87c4cb33-338c-4949-aace-b5012e1977c6') " }
MongoDB server version: 4.2.2
Server has startup warnings:
2020-10-01T16:00:08.536+0530 I CONTROL [initandlisten]
2020-10-01T16:00:08.536+0530 I CONTROL [initandlisten] ** WARNING: Access control is not enabled for the database.
2020-10-01T16:00:08.537+0530 I CONTROL [initandlisten] **           Read and write access to data and configuration is unrestricted.
2020-10-01T16:00:08.586+0530 I CONTROL [initandlisten]

Enable MongoDB's free cloud-based monitoring service, which will then receive and display
metrics about your deployment (disk utilization, CPU, operation statistics, etc).

The monitoring data will be available on a MongoDB website with a unique URL accessible to you
and anyone you share the URL with. MongoDB may use this information to make product
improvements and to suggest MongoDB products and deployment options to you.

To enable free monitoring, run the following command: db.enableFreeMonitoring()
To permanently disable this reminder, run the following command: db.disableFreeMonitoring()
---
> use mynew;
switched to db mynew
> db.Customer.find();
[{"_id": "ObjectId('5f7f289fa8939e106c990cc2')", "cust_id": 1, "Acc_bal": 2500, "Acc_type": "Z"}, {"_id": "ObjectId('5f7f289fa8939e106c990cc3')", "cust_id": 2, "Acc_bal": 3000, "Acc_type": "A"}, {"_id": "ObjectId('5f7f289fa8939e106c990cc4')", "cust_id": 3, "Acc_bal": 1200, "Acc_type": "A"}, {"_id": "ObjectId('5f7f289fa8939e106c990cc5')", "cust_id": 4, "Acc_bal": 2400, "Acc_type": "A"}, {"_id": "ObjectId('5f7f289fa8939e106c990cc6')", "cust_id": 5, "Acc_bal": 1200, "Acc_type": "Z"}, {"_id": "ObjectId('5f7f289fa8939e106c990cc7')", "cust_id": 6, "Acc_bal": 1200, "Acc_type": "Z"}]

```

Q3.)

Perform the following DB operations using Cassandra.

- A. Create a keyspace by name Employee
- B. Create a column family by name Employee-Info with attributes Emp_Id Primary Key, Emp_Name, Designation, Date_of_Joining, Salary, Dept_Name.
- C. Insert the values into the table in batch
- D. Update Employee name and Department of Emp-Id 121.
- E. Sort the details of Employee records based on salary.
- F. Alter the schema of the table Employee_Info to add a column Projects which stores a set of Projects done by the corresponding Employee.
- G. Update the altered table to add project names.
- H. Create a TTL of 15 seconds to display the values of Employees.

CODE – 3

```
cqlsh> CREATE KEYSPACE Employees WITH replication =  
{'class':'SimpleStrategy','replication_factor':3};
```

```
cqlsh> use employees;
```

```
cqlsh:employees> CREATE COLUMNFAMILY employee_info(Emp_id INT PRIMARY KEY,  
Emp_name VARCHAR, Designation VARCHAR, Date_of_joining VARCHAR, Salary FLOAT,  
Dept_name VARCHAR);
```

```
cqlsh:employees> BEGIN BATCH INSERT INTO employee_info (emp_id , emp_name ,  
designation , date_of_joining , salary , dept_name ) values  
(119,'Shivam','CEO','01/06/2020',10000, 'IT');
```

```
... INSERT INTO employee_info (emp_id , emp_name , designation , date_of_joining ,  
salary , dept_name ) values (120,'Shivam Raj', 'CFO','01/01/2020',5000,'Finance');
```

```
... INSERT INTO employee_info (emp_id , emp_name , designation , date_of_joining ,  
salary , dept_name ) values (121,'Shiv', 'HR head','01/04/2020',4000,'SALES');
```

```
... INSERT INTO employee_info (emp_id , emp_name , designation , date_of_joining ,  
salary , dept_name ) values (122,'Raj','CTO','01/06/2020',10000, 'SALES'); APPLY BATCH;
```

```
cqlsh:employees> SELECT * FROM employee_info;
```

```
cqlsh:employees> UPDATE employee_info SET emp_name = 'Shiva',dept_name = 'IT' WHERE emp_id = 121;
```

```
cqlsh:employees> SELECT * FROM employee_info;
```

```
cqlsh:employees> ALTER TABLE employee_info ADD PROJECT SET<VARCHAR>;
```

```
cqlsh:employees> UPDATE employee_info SET project=project+{'Covid-19','BDA Analysis'} WHERE emp_id=122;
```

```
cqlsh:employees> INSERT INTO employee_info (emp_id , emp_name , designation , date_of_joining , salary , dept_name ) values (123,'Random Guy','Manager','01/06/2020',10000, 'IT') USING TTL 15;
```

```
cqlsh:employees> SELECT TTL(designation) FROM employee_info where emp_id=123;
```

```
cqlsh:employees> SELECT * FROM employee_info;
```

```
cqlsh:employees> SELECT TTL(designation) FROM employee_info where emp_id=123;
```

Write-up-3

Shivam Raj BDA Lab
IBM17C3095 B3
Lab - 6 (I)

classmate

Date

Page

> cqlsh

```
CREATE KEYSPACE Employees WITH replication =  
{ 'class': 'SimpleStrategy', 'replication-factor': 3 };  
use employees;
```

```
CREATE COLUMNFAMILY employee-info (emp-id  
INT PRIMARY KEY, emp-name VARCHAR,  
Designation VARCHAR, Date-of-joining VARCHAR,  
Salary FLOAT, Dept-name VARCHAR);
```

```
BEGIN BATCH INSERT INTO employee-info (emp-  
id, emp-name, designation, date-of-joining,  
Salary, dept-name) VALUES (11, 'Shivam', 'CEO',  
'01/06/2020', 10000, 'IT');
```

...

```
... APPLY BATCH;
```

```
SELECT * FROM employee-info;
```

```
UPDATE employee-info SET emp-name = 'Shiva',  
dept-name = 'HR' WHERE emp-id = 121;
```

```
ALTER TABLE employee-info ADD PROJECT  
SET<VARCHAR>;
```

```
UPDATE employee-info SET project = project + 1  
'COVID-19', 'BDA Analysis' WHERE emp-id = 122;  
SELECT * FROM employee-info
```

```
BEGIN BATCH INSERT INTO employee-info (emp-id,  
emp-name, designation, date-of-joining, salary, dept-name)  
VALUES (123, 'Random Guy', 'Manager', '01/06/2020', 10000, 'IT')  
USING TTL 15; ... APPLY BATCH;
```

```
SELECT TTL(designation) FROM employee-info WHERE emp-id=123;
```


Screenshot – 3

```
Administrator: Command Prompt - cqlsh
Microsoft Windows [Version 10.0.18363.1139]
(c) 2019 Microsoft Corporation. All rights reserved.

C:\WINDOWS\system32>cd C:\Program Files\apache-cassandra-3.11.8\bin
C:\Program Files\apache-cassandra-3.11.8\bin>cqlsh

WARNING: console codepage must be set to cp65001 to support utf-8 encoding on Windows platforms.
If you experience encoding problems, change your console codepage with 'chcp 65001' before starting cqlsh.

Connected to Test Cluster at 127.0.0.1:9042.
[cqlsh 5.0.1 | Cassandra 3.11.8 | CQL spec 3.4.4 | Native protocol v4]
Use HELP for help.
WARNING: pyreadline dependency missing. Install to enable tab completion.
cqlsh> CREATE KEYSPACE Employees WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 3};
cqlsh> use employees;
cqlsh:employees> CREATE COLUMNFAMILY employee_info (emp_id INT PRIMARY KEY, emp_name VARCHAR, Designation VARCHAR, Date_of_joining VARCHAR, Salary FLOAT, Dept_name VARCHAR);
cqlsh:employees> BEGIN BATCH INSERT INTO employee_info (emp_id, emp_name, designation, date_of_joining, salary, dept_name) values (11, 'Shivam', 'CEO', '01/06/2020', 10000, 'IT');
... INSERT INTO employee_info (emp_id, emp_name, designation, date_of_joining, salary, dept_name) values (12, 'Shivam Raj', 'CTO', '01/06/2020', 10000, 'Finance');
... INSERT INTO employee_info (emp_id, emp_name, designation, date_of_joining, salary, dept_name) values (121, 'Shiva', 'HR head', '01/04/2020', 4000, 'SALES');
... INSERT INTO employee_info (emp_id, emp_name, designation, date_of_joining, salary, dept_name) values (122, 'Raj', 'CTO', '01/06/2020', 10000, 'SALES');
... APPLY BATCH;
cqlsh:employees> select * from employee_info;

emp_id | date_of_joining | dept_name | designation | emp_name | salary
-----|-----|-----|-----|-----|-----
11 | 01/06/2020 | IT | CEO | Shivam | 10000
122 | 01/06/2020 | SALES | CTO | Raj | 10000
121 | 01/04/2020 | SALES | HR head | Shiva | 4000
12 | 01/01/2020 | Finance | CFO | Shivam Raj | 5000
(4 rows)
cqlsh:employees> UPDATE employee_info SET emp_name = 'Shiva', dept_name = 'HR' WHERE emp_id = 121;
cqlsh:employees> select * from employee_info;

emp_id | date_of_joining | dept_name | designation | emp_name | salary
-----|-----|-----|-----|-----|-----
11 | 01/06/2020 | IT | CEO | Shivam | 10000
122 | 01/06/2020 | SALES | CTO | Raj | 10000
121 | 01/04/2020 | HR | HR head | Shiva | 4000
12 | 01/01/2020 | Finance | CFO | Shivam Raj | 5000
(4 rows)
cqlsh:employees> ALTER TABLE employee_info ADD PROJECT SET<VARCHAR>;
cqlsh:employees> select * from employee_info;

emp_id | date_of_joining | dept_name | designation | emp_name | project | salary
-----|-----|-----|-----|-----|-----|-----
11 | 01/06/2020 | IT | CEO | Shivam | null | 10000
122 | 01/06/2020 | SALES | CTO | Raj | null | 10000
121 | 01/04/2020 | HR | HR head | Shiva | null | 4000
12 | 01/01/2020 | Finance | CFO | Shivam Raj | null | 5000
(4 rows)
cqlsh:employees> UPDATE employee_info SET project=project+'Covid-19','BDA Analysis' WHERE emp_id=122;
cqlsh:employees> select * from employee_info;

emp_id | date_of_joining | dept_name | designation | emp_name | project | salary
-----|-----|-----|-----|-----|-----|-----
11 | 01/06/2020 | IT | CEO | Shivam | null | 10000
122 | 01/06/2020 | SALES | CTO | Raj | ('BDA Analysis', 'Covid-19') | 10000
121 | 01/04/2020 | HR | HR head | Shiva | null | 4000
12 | 01/01/2020 | Finance | CFO | Shivam Raj | null | 5000
(4 rows)
cqlsh:employees> BEGIN BATCH INSERT INTO employee_info (emp_id, emp_name, designation, date_of_joining, salary, dept_name) values (123, 'Random Guy', 'Manager', '01/06/2020', 10000, 'IT') USING TTL 15;
... APPLY BATCH;
cqlsh:employees> select * from employee_info;

emp_id | date_of_joining | dept_name | designation | emp_name | project | salary
-----|-----|-----|-----|-----|-----|-----
11 | 01/06/2020 | IT | CEO | Shivam | null | 10000
122 | 01/06/2020 | SALES | CTO | Raj | ('BDA Analysis', 'Covid-19') | 10000
121 | 01/04/2020 | HR | HR head | Shiva | null | 4000
12 | 01/01/2020 | Finance | CFO | Shivam Raj | null | 5000
(4 rows)
cqlsh:employees> SELECT TTL(designation) FROM employee_info where emp_id=123;

ttl(designation)
-----
(0 rows)
cqlsh:employees> BEGIN BATCH INSERT INTO employee_info (emp_id, emp_name, designation, date_of_joining, salary, dept_name) values (123, 'Random Guy', 'Manager', '01/06/2020', 10000, 'IT') USING TTL 15;
... APPLY BATCH;
cqlsh:employees> SELECT TTL(designation) FROM employee_info where emp_id=123;

ttl(designation)
-----
5
(1 rows)
cqlsh:employees>
```

Q4.)

Perform the following DB operations using Cassandra.

- A. Create a keyspace by name Library.
- B. Create a column family by name Library-Info with attributes Stud_Id Primary Key, Counter_value of type Counter, Stud_Name, Book-Name, Book-Id, Date_of_issue.
- C. Insert the values into the table in batch.
- D. Display the details of the table created and increase the value of the counter.
- E. Write a query to show that a student with id 112 has taken a book "BDA" 2 times.
- F. Export the created column to a csv file.
- G. Import a given csv dataset from local file system into Cassandra column family

CODE – 4

```
cqlsh> CREATE KEYSPACE Library WITH replication =  
{'class':'SimpleStrategy','replication_factor':3};
```

```
cqlsh> use Library;
```

```
cqlsh:library> CREATE COLUMNFAMILY library_info(stud_id int, counter_value counter,  
stud_name VARCHAR, book_name VARCHAR, book_id INT, DOI VARCHAR, PRIMARY  
KEY(stud_id,stud_name,book_name,book_id,doi));
```

```
cqlsh:library> UPDATE library_info set counter_value=counter_value+1 where stud_id=112  
and stud_name='Shivam' and book_name='Vedas' and book_id=1 and doi='01/01/2020';
```

```
cqlsh:library> UPDATE library_info set counter_value=counter_value+1 where stud_id=112  
and stud_name='Shivam' and book_name='ved' and book_id=1 and doi='01/01/2020';
```

```
cqlsh:library> UPDATE library_info set counter_value=counter_value+1 where stud_id=113  
and stud_name='Raj' and book_name='Geeta' and book_id=2 and doi='01/05/2020';
```

```
cqlsh:library> SELECT * FROM library_info;
```

```
cqlsh:library> SELECT * FROM library_info WHERE stud_id=112;
```

```
cqlsh:library> SELECT book_name,counter_value FROM library_info WHERE stud_id=112;
```

```
cqlsh:library> COPY library_info (stud_id , counter_value , stud_name , book_name ,  
book_id , doi ) TO 'library.csv' WITH HEADER=TRUE;
```

```
cqlsh:library> truncate library_info ;
```

```
cqlsh:library> SELECT * FROM library_info;
```

```
cqlsh:library> COPY library_info (stud_id , counter_value , stud_name , book_name ,  
book_id , doi ) FROM 'library.csv' WITH HEADER=TRUE;
```

```
cqlsh:library> SELECT * FROM library_info;
```

Write-up-4

Shivam Raj
IBM17CS095

BDA Lab
B3

Lab - 6 (II)

classmate

Date _____
Page _____

> cqlsh

```
CREATE KEYSPACE Lib WITH replication =  
{ 'class': 'SimpleStrategy', 'replication-factor': 3 };  
use Lib;
```

```
CREATE COLUMNFAMILY library-info (stud-id int,  
counter_value counter, stud-name VARCHAR, book-name  
VARCHAR, book-id INT, doi VARCHAR, PRIMARY KEY  
(stud-id, stud-name, book-name, book-id, doi));
```

```
UPDATE library-info set counter_value =  
counter_value + 1 where stud-id = 11 and  
stud-name = 'Shivam' and book-name =  
'Vedas' and book-id = 1 and doi = '01/01/2020';
```

```
Select * from library-info where stud-id = 11;
```

```
COPY library-info (stud-id, counter_value, stud-name,  
book-id, doi) TO 'D:\BDA\cassandra\library.csv'  
WITH HEADER = TRUE;
```

```
COPY library-info (stud-id, counter_value, stud-name,  
book-name, book-id, doi) FROM 'D:\BDA\  
cassandra\library.csv' WITH HEADER = TRUE;
```

```
Select * from library-info;
```

Screenshot – 4

```
Command Prompt - cqlsh
Microsoft Windows [Version 10.0.18363.1139]
(c) 2019 Microsoft Corporation. All rights reserved.

C:\Users\SHIVAM>cqlsh

WARNING: console codepage must be set to cp65001 to support utf-8 encoding on Windows platforms.
If you experience encoding problems, change your console codepage with 'chcp 65001' before starting cqlsh.

Connected to Test Cluster at 127.0.0.1:9042.
[cqlsh 5.0.1 | Cassandra 3.11.8 | CQL spec 3.4.4 | Native protocol v4]
Use HELP for help.
WARNING: pyreadline dependency missing. Install to enable tab completion.
cqlsh> CREATE KEYSPACE lib WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 3};
cqlsh> use lib;
cqlsh:lib> CREATE COLUMNFAMILY library_info(stud_id int, counter_value counter, stud_name VARCHAR, book_name VARCHAR, book_id INT, doi VARCHAR, PRIMARY KEY(stud_id,stud_name,book_name,book_id,doi));
cqlsh:lib> UPDATE library_info set counter_value=counter_value+1 where stud_id=11 and stud_name='Shivam' and book_name='Vedas' and book_id=1 and doi='01/01/2020';
cqlsh:lib> UPDATE library_info set counter_value=counter_value+1 where stud_id=11 and stud_name='Shivam' and book_name='Vedas' and book_id=1 and doi='01/01/2020';
cqlsh:lib> UPDATE library_info set counter_value=counter_value+1 where stud_id=10 and stud_name='Raj' and book_name='Geeta' and book_id=2 and doi='01/05/2020';
cqlsh:lib> SELECT * FROM library_info;

stud_id | stud_name | book_name | book_id | doi | counter_value
-----|-----|-----|-----|-----|-----
10 | Raj | Geeta | 2 | 01/05/2020 | 1
11 | Shivam | Vedas | 1 | 01/01/2020 | 2
(2 rows)
cqlsh:lib> SELECT * FROM library_info WHERE stud_id=11;

stud_id | stud_name | book_name | book_id | doi | counter_value
-----|-----|-----|-----|-----|-----
11 | Shivam | Vedas | 1 | 01/01/2020 | 2
(1 rows)
cqlsh:lib> SELECT * from library_info WHERE counter_value=2 ALLOW FILTERING;

stud_id | stud_name | book_name | book_id | doi | counter_value
-----|-----|-----|-----|-----|-----
11 | Shivam | Vedas | 1 | 01/01/2020 | 2
(1 rows)
cqlsh:lib> SELECT book_name,counter_value FROM library_info WHERE stud_id=11;

book_name | counter_value
-----|-----
Vedas | 2
(1 rows)
cqlsh:lib> COPY library_info (stud_id , counter_value , stud_name , book_name , book_id , doi ) TO 'D:\BDA\Cassandra\library.csv' WITH HEADER=TRUE;
Using 3 child processes

Command Prompt - cqlsh
cqlsh:lib> SELECT book_name,counter_value FROM library_info WHERE stud_id=11;

book_name | counter_value
-----|-----
Vedas | 2
(1 rows)
cqlsh:lib> COPY library_info (stud_id , counter_value , stud_name , book_name , book_id , doi ) TO 'D:\BDA\Cassandra\library.csv' WITH HEADER=TRUE;
Using 3 child processes

Starting copy of lib.library_info with columns [stud_id, counter_value, stud_name, book_name, book_id, doi].
Processed: 2 rows; Rate: 1 rows/s; Avg. rate: 1 rows/s
2 rows exported to 1 files in 3.140 seconds.
cqlsh:lib> truncate library_info ;
cqlsh:lib> SELECT * FROM library_info;

stud_id | stud_name | book_name | book_id | doi | counter_value
-----|-----|-----|-----|-----|-----
(0 rows)
cqlsh:lib> COPY library_info (stud_id , counter_value , stud_name , book_name , book_id , doi ) FROM 'D:\BDA\Cassandra\library.csv' WITH HEADER=TRUE;
Using 3 child processes

Starting copy of lib.library_info with columns [stud_id, counter_value, stud_name, book_name, book_id, doi].
Process ImportProcess-4: 1 rows/s; Avg. rate: 1 rows/s
Process ImportProcess-5:
Traceback (most recent call last):
  File "C:\Python27\lib\multiprocessing\process.py", line 267, in _bootstrap
    self.run()
  File "C:\Python27\lib\multiprocessing\process.py", line 267, in _bootstrap
    self.run()
  File "C:\Program Files\apache-cassandra-3.11.8\bin\..pylib\cqlshlib\copyutil.py", line 2328, in run
  File "C:\Program Files\apache-cassandra-3.11.8\bin\..pylib\cqlshlib\copyutil.py", line 2328, in run
    self.close()
  File "C:\Program Files\apache-cassandra-3.11.8\bin\..pylib\cqlshlib\copyutil.py", line 2332, in close
  File "C:\Program Files\apache-cassandra-3.11.8\bin\..pylib\cqlshlib\copyutil.py", line 2332, in close
    self.session.cluster.shutdown()
  File "C:\Program Files\apache-cassandra-3.11.8\bin\..pylib\cassandra-driver-internal-only-3.11.0-bb96859b.zip\cassandra-driver-3.11.0-bb96859b\cassandra\cluster.py", line 1259, in shutdown
  File "C:\Program Files\apache-cassandra-3.11.8\bin\..pylib\cassandra-driver-internal-only-3.11.0-bb96859b.zip\cassandra-driver-3.11.0-bb96859b\cassandra\cluster.py", line 1259, in shutdown
    self.control_connection.shutdown()
  File "C:\Program Files\apache-cassandra-3.11.8\bin\..pylib\cassandra-driver-internal-only-3.11.0-bb96859b.zip\cassandra-driver-3.11.0-bb96859b\cassandra\cluster.py", line 2850, in shutdown
  File "C:\Program Files\apache-cassandra-3.11.8\bin\..pylib\cassandra-driver-internal-only-3.11.0-bb96859b.zip\cassandra-driver-3.11.0-bb96859b\cassandra\cluster.py", line 2850, in shutdown
    self.connection.close()
  File "C:\Program Files\apache-cassandra-3.11.8\bin\..pylib\cassandra-driver-internal-only-3.11.0-bb96859b.zip\cassandra-driver-3.11.0-bb96859b\cassandra\io\asyncoreactor.py", line 373, in close
  File "C:\Program Files\apache-cassandra-3.11.8\bin\..pylib\cassandra-driver-internal-only-3.11.0-bb96859b.zip\cassandra-driver-3.11.0-bb96859b\cassandra\io\asyncoreactor.py", line 373, in close
```


Q5.)

Develop a MapReduce program to count the number of occurrences of words in a given file.

CODE – 5

```
cd /usr/local/hadoop/sbin
```

```
start-all.sh
```

```
jps
```

```
mkdir hadfiles
```

```
cd hadfiles
```

```
touch wordcountfile.txt
```

```
nano wordcountfile.txt
```

```
hadoop fs -mkdir/programs
```

```
hadoop fs -copyFromLocal/home/hadoop/hadfiles/wordcountfiles.txt/programs/prog1.txt
```

```
hadoop jar/home/hadoop/hadfiles/wordcount.jar
```

```
WordCount/programs/prog1.txt/programs/output
```

```
hadoop fs -ls/programs/output
```

```
hadoop fs -cat/programs/output2/part-r-00000
```

```
stop-all.sh
```

Screenshot – 5

1. Starting Hadoop Cluster

```
[/]$ cd /usr/local/hadoop/sbin
[/usr/local/hadoop/sbin]$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [DESKTOP-72U025C]
2020-12-08 23:10:37,345 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Starting resourcemanager
Starting nodemanagers
[/usr/local/hadoop/sbin]$ jps
14546 NameNode
15315 NodeManager
14948 SecondaryNameNode
15159 ResourceManager
15671 Jps
14698 DataNode
```

2. Creating file for counting words of

```
[~]$ mkdir hadfiles
[~]$ cd hadfiles
[~/hadfiles]$ touch wordcountfile.txt
[~/hadfiles]$ nano wordcountfile.txt
```

3. Moving file to Hadoop File System

```
[/usr/local/hadoop]$ hadoop fs -mkdir /programs
```

```
[/usr/local/hadoop]$ hadoop fs -copyFromLocal /home/hadoop/hadfiles/wordcountfile.txt /programs/prog1.txt
```

4. Running wordcount.jar, which consists of wordcount\$Map, wordcount\$Reduce, and wordcount classes

```
[~/hadfiles]$ hadoop jar /home/hadoop/hadfiles/wordcount.jar WordCount /programs/prog1.txt /programs/output/
```

5. Output received

```
~/hadfiles
❶ hadoop fs -ls /programs/output
Found 2 items
-rw-r--r-- 1 hdoop supergroup          0 2020-12-08 23:36 /programs/output/_SUCCESS
-rw-r--r-- 1 hdoop supergroup      69 2020-12-08 23:36 /programs/output/part-r-000000

~/hadfiles
❶ hadoop fs -cat /programs/output/part-r-000000
are 1
brother 1
family 1
hi 1
how 5
is 4
job 1
sister 1
you 1
your 4
```

6. Stopping Hadoop

```
~/hadfiles
❶ stop-all.sh
```

Q6.)

For the given file, create a Map Reduce program to

- a) Find the average temperature for each year from NCDC data set.

CODE – 6

```
cd /usr/local/hadoop/sbin
```

```
start-all.sh
```

```
jps
```

```
hadoop fs -copyFromLocal/home/hadfiles/1901/programs/avgtemp.txt
```

```
hadoop fs-ls/programs
```

```
hadoop jar/home/hadoop/hadfiles/average.jar AverageDriver/programs/avgtemp.txt/progr
```

```
hadoop fs -ls/programs
```

```
hadoop fs -cat/programs/output2/part-r-00000
```

```
stop-all.sh
```


Screenshot – 6

1. Starting Hadoop Cluster

```
hadoop@DESKTOP-72U025C /$ cd /usr/local/hadoop/sbin
hadoop@DESKTOP-72U025C /usr/local/hadoop/sbin$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [DESKTOP-72U025C]
Starting resourcemanager
Starting nodemanagers
hadoop@DESKTOP-72U025C /usr/local/hadoop/sbin$ jps
4128 SecondaryNameNode
3906 DataNode
4499 NodeManager
4342 ResourceManager
4859 Jps
3725 NameNode
```

2. Copying Binary file to Hadoop File System as a txt file

```
hadoop@DESKTOP-72U025C ~/hadfiles$ hadoop fs -copyFromLocal /home/hadoop/hadfiles/1901 /programs/avgtemp.txt
hadoop@DESKTOP-72U025C ~/hadfiles$ hadoop fs -ls /programs
Found 3 items
-rw-r--r-- 1 hadoop supergroup 888190 2020-12-14 15:04 /programs/avgtemp.txt
```

3. Running Average.jar

```
hadoop@DESKTOP-72U025C ~/hadfiles$ hadoop jar /home/hadoop/hadfiles/average.jar AverageDriver /programs/avgtemp.txt /programs/output2
```

4. Checking location of output file

```
hadoop@DESKTOP-72U025C ~/hadfiles$ hadoop fs -ls /programs
Found 4 items
-rw-r--r-- 1 hadoop supergroup 888190 2020-12-14 15:04 /programs/avgtemp.txt
drwxr-xr-x - hadoop supergroup 0 2020-12-08 23:36 /programs/output
drwxr-xr-x - hadoop supergroup 0 2020-12-14 15:07 /programs/output2
-rw-r--r-- 1 hadoop supergroup 92 2020-12-08 23:20 /programs/prog1.txt
```

5. Result

```
hadoop@DESKTOP-72U025C ~/hadfiles$ hadoop fs -cat /programs/output2/part-r-00000
1901 46
```

6. Stopping Hadoop

```
hadoop@DESKTOP-72U025C ~/hadfiles$ stop-all.sh
```

Q7.)

Write Queries in Hive to do the following

- A. Create an external table named with the following attributes -> Empl_ID -> Emp_Name -> Designation -> Salary
- B. Load data into table from a given file.
- C. Create a view to Generate a query to retrieve the employee details who earn a salary of more than Rs 30000.
- D. Alter the table to add a column Dept_Id and Generate a query to retrieve the employee details in order by using Dept_Id.
- E. Generate a query to retrieve the number of employees in each department whose salary is greater than 30000
- F. Create another table Department with attributes -> Dept_Id -> Dept_name -> Emp_Id 7. Display the cumulative details of each employee along with department details

CODE – 7

A.

```
CREATE DATABASE IF NOT EXISTS lab9 COMMENT 'employee program' WITH DBPROPERTIES
('creator'=ROUNAK);
SHOW DATABASES;
DESCRIBE DATABASE lab9;
USE lab9;
CREATE EXTERNAL TABLE IF NOT EXISTS Employee(EmplID INT,EmpName
STRING,Designation STRING,Salary FLOAT) ROW FORMAT DELIMITED FIELDS TERMINATED
BY '\t';
```

B.

```
LOAD DATA LOCAL INPATH '/home/rounak/Desktop/employeeInput.txt' OVERWRITE INTO
```

```
TABLE Employee;  
SELECT * FROM Employee;
```

C.

```
CREATE VIEW emp_30000 AS SELECT * FROM Employee WHERE Salary>30000;  
SELECT * FROM emp_30000;
```

D.

```
ALTER TABLE Employee ADD COLUMNS(DeptID INT);  
LOAD DATA LOCAL INPATH '/home/rounak/Desktop/employeeInputAltered.txt'  
OVERWRITE INTO TABLE Employee;  
SELECT * FROM Employee;  
SELECT * FROM Employee ORDER BY DeptID;
```

E.

```
SELECT DeptID,count(*) FROM Employee WHERE Salary>=30000 GROUP BY DeptID;
```

F.

```
CREATE EXTERNAL TABLE IF NOT EXISTS Department(DeptId INT,DeptName STRING) ROW  
FORMAT DELIMITED FIELDS TERMINATED BY '\t';  
LOAD DATA LOCAL INPATH '/home/rounak/Desktop/DepartmentInput.txt' OVERWRITE  
INTO TABLE Department;  
SELECT * FROM Department;
```

G.

```
SELECT a.EmpID,a.EmpName,a.Designation,a.Salary,b.DeptName FROM Employee a  
JOIN Department b ON a.DeptID=b.DeptId;
```

Screenshot – 7

```
Activities Terminal Sat 03:15
hduser@ubuntu: /usr/local/apache-hive-2.1.0-bin/bin

hive> SELECT * FROM Employee ORDER BY DeptID;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = hduser_20201219031430_f37e5e18-d55b-4d50-8836-b3d86a9ca242
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Job running in-process (local Hadoop)
2020-12-19 03:14:36,260 Stage-1 Map = 0%, reduce = 0%
2020-12-19 03:14:45,727 Stage-1 map = 100%, reduce = 100%
Ended Job = job_local426704584_0001
MapReduce Jobs Launched:
Stage-Stage-1: HDFS Read: 1944 HDFS Write: 972 SUCCESS
Total MapReduce CPU Time Spent: 0 msec
OK
1      Kelly      Sr. Manager      60000.0 11
2      John       Manager          50000.0 11
7      Rohan      HR manager       25000.0 12
8      Joana      Software Engineer 23000.0 13
4      Kedrik     Software Engineer 30000.0 13
3      Harry      Software Engineer 45000.0 13
6      Ubern      Test Engineer    20000.0 15
5      Tom        Test Engineer    20000.0 15
Time taken: 15.476 seconds, Fetched: 8 row(s)
```

```
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Job running in-process (local Hadoop)
2020-12-19 03:18:13,778 Stage-3 map = 100%, reduce = 0%
Ended Job = job_local1327814845_0003
MapReduce Jobs Launched:
Stage-Stage-3: HDFS Read: 1542 HDFS Write: 546 SUCCESS
Total MapReduce CPU Time Spent: 0 msec
OK
1      John       Manager          50000.0 Business Management
2      Kelly      Sr. Manager      60000.0 Business Management
3      Harry      Software Engineer 45000.0 Development
4      Kedrik     Software Engineer 30000.0 Development
5      Tom        Test Engineer    20000.0 Testing
6      Ubern      Test Engineer    20000.0 Testing
7      Rohan      HR manager       25000.0 HR
8      Joana      Software Engineer 23000.0 Development
Time taken: 51.043 seconds, Fetched: 8 row(s)
```

```
Activities Terminal Sat 03:08
hduser@ubuntu: /usr/local/apache-hive-2.1.0-bin/bin

File Edit View Search Terminal Help
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Logging initialized using configuration in jar:file:/usr/local/apache-hive-2.1.0-bin/lib/hive-common-2.1.0.jar!/hive-log4j2.
properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution eng
ine (i.e. spark, tez) or using Hive 1.X releases.
hive> CREATE DATABASE IF NOT EXISTS lab9 COMMENT 'employee program' WITH DBPROPERTIES ('creator'='CHANDANA');
OK
Time taken: 3.886 seconds
hive> SHOW DATABASES;
OK
default
lab9
Time taken: 1.768 seconds, Fetched: 2 row(s)
hive> DESCRIBE DATABASE lab9;
OK
lab9    hive employee program    hdfs://localhost:54310/user/hive/warehouse/lab9.db    hduser    USER
Time taken: 0.104 seconds, Fetched: 1 row(s)
hive> USE lab9;
OK
Time taken: 0.08 seconds
hive> CREATE EXTERNAL TABLE IF NOT EXISTS Employee(EmpID INT,EmpName STRING,Designation STRING,Salary FLOAT) ROW FORMAT DELI
MITED FIELDS TERMINATED BY '\t';
OK
Time taken: 1.98 seconds
hive> LOAD DATA LOCAL INPATH '/home/chandana/Desktop/employeeInput.txt' OVERWRITE INTO TABLE Employee;
Loading data to table lab9.employee
OK
Time taken: 5.339 seconds
hive>
```